

B529: Homework 3

Nathan Byers

Wednesday, April 7, 2015

Question 1

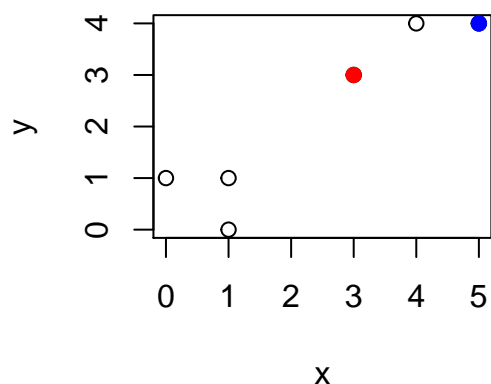
In the clustering problems, the input data set contains the following six data points (1,1), (0,1), (1,0), (3,3) (4,4), (5,4). Manhattan distance is chosen to compute the distance between two data points. The k-means clustering method is used to find the clusters with $k = 2$ and initial centroids (3,3) and (5,4). Please describe the steps of the k-means clustering. (In each step, provide the information about cluster assignments of data points, centroids; and when the algorithm stops). (25 points)

Answer 1

First I read the points into R and plot them.

```
x <- c(1, 0, 1, 3, 4, 5)
y <- c(1, 1, 0, 3, 4, 4)
seed1 <- c(3, 3)
seed2 <- c(5, 4)

plot(x = x, y = y)
points(c(seed1[1], seed2[1]),
       c(seed1[2], seed2[2]),
       col = c("red", "blue"),
       pch = 19)
```



Then I create a data frame that gives the initial state of the points.

```

group.df <- data.frame(x, y, group = NA,
                      centr_x = NA, centr_y = NA)
group.df[group.df$x == 3 &
          group.df$y == 3,
          c("group", "centr_x", "centr_y")] <- c(1, 3, 3)

group.df[group.df$x == 5 &
          group.df$y == 4,
          c("group", "centr_x", "centr_y")] <- c(2, 5, 4)

manDist <- function(x, y, centr_x, centr_y){
  abs(centr_x - x) +
  abs(centr_y - y)
}

group.df$dist1 <- unlist(mapply(manDist, group.df$x,
                               group.df$y,
                               MoreArgs = list(
                                 centr_x = seed1[1],
                                 centr_y = seed1[2])))

group.df$dist2 <- unlist(mapply(manDist, group.df$x,
                               group.df$y,
                               MoreArgs = list(
                                 centr_x = seed2[1],
                                 centr_y = seed2[2])))

group.df$group <- apply(group.df, 1, function(row){
  if(row["dist1"] ==
      min(row[c("dist1", "dist2"))]){
    1
  }else{2}
})

group.df$change <- NA

group.df

```

```

##   x y group centr_x centr_y dist1 dist2 change
## 1 1 1     1     NA     NA     4     7     NA
## 2 0 1     1     NA     NA     5     8     NA
## 3 1 0     1     NA     NA     5     8     NA
## 4 3 3     1      3      3     0     3     NA
## 5 4 4     2     NA     NA     2     1     NA
## 6 5 4     2      5      4     3     0     NA

```

Now we find the new centroids, recalculate the distances from each point to the new centroids, and regroup the points according to the shortest Manhattan distance to the new centroids.

```

reGroup <- function(df){
  center_x1 <- mean(df[df$group == 1, "x"])
  center_y1 <- mean(df[df$group == 1, "y"])
  center_x2 <- mean(df[df$group == 2, "x"])

```

```

center_y2 <- mean(df[df$group == 2, "y"])
df[df$group == 1, "centr_x"] <- center_x1
df[df$group == 1, "centr_y"] <- center_y1
df[df$group == 2, "centr_x"] <- center_x2
df[df$group == 2, "centr_y"] <- center_y2
df$dist1 <- unlist(mapply(manDist, df$x,
                        df$y,
                        MoreArgs = list(
                          centr_x = center_x1,
                          centr_y = center_y1)))
df$dist2 <- unlist(mapply(manDist, df$x,
                        df$y,
                        MoreArgs = list(
                          centr_x = center_x2,
                          centr_y = center_y2)))

g.initial <- df$group
g <- apply(df, 1, function(row){
  if(row["dist1"] != row["dist2"]){
    if(row["dist1"] ==
        min(row[c("dist1", "dist2")])){
      1
    }else{2}
  }else{0}
})
df$group <- g
df$change <- sapply(1:length(g), function(i){
  !(g[i] == g.initial[i])
})
df
}

group.df2 <- reGroup(group.df)

group.df2

```

```

##   x y group centr_x centr_y dist1 dist2 change
## 1 1 1     1   1.25   1.25   0.5   6.5  FALSE
## 2 0 1     1   1.25   1.25   1.5   7.5  FALSE
## 3 1 0     1   1.25   1.25   1.5   7.5  FALSE
## 4 3 3     2   1.25   1.25   3.5   2.5   TRUE
## 5 4 4     2   4.50   4.00   5.5   0.5  FALSE
## 6 5 4     2   4.50   4.00   6.5   0.5  FALSE

```

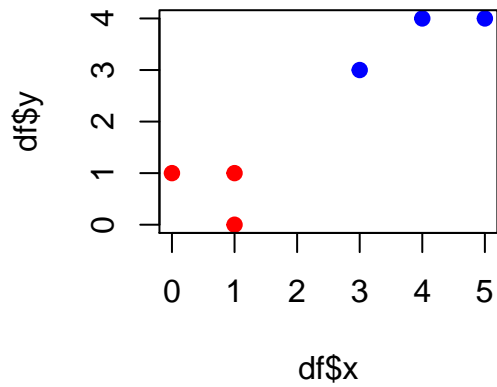
We can see that at least one point changed. Below is a plot of the new groups.

```

plotGroups <- function(df){
  group.cols <- df$group
  group.cols[group.cols == 1] <- "red"
  group.cols[group.cols == 2] <- "blue"
  group.cols[group.cols == 0] <- "black"
  plot(df$x, df$y, pch = 19,
        col = group.cols)
}

```

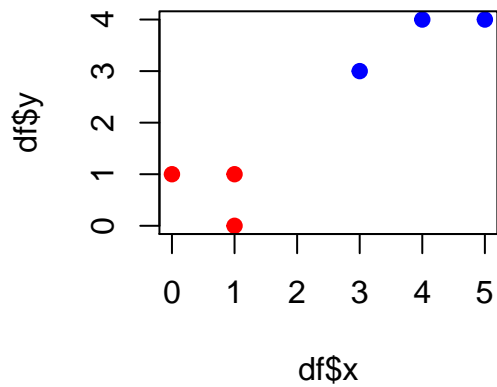
```
plotGroups(group.df2)
```



Now we regroup again, because at least one point changed groups in the last iteration.

```
group.df3 <- reGroup(group.df2)
```

```
plotGroups(group.df3)
```



```
group.df3
```

```
##   x y group  centr_x  centr_y  dist1  dist2 change
## 1 1 1     1 0.6666667 0.6666667 0.6666667 5.6666667 FALSE
## 2 0 1     1 0.6666667 0.6666667 1.0000000 6.6666667 FALSE
## 3 1 0     1 0.6666667 0.6666667 1.0000000 6.6666667 FALSE
```

```
## 4 3 3      2 4.0000000 3.6666667 4.6666667 1.6666667 FALSE
## 5 4 4      2 4.0000000 3.6666667 6.6666667 0.3333333 FALSE
## 6 5 4      2 4.0000000 3.6666667 7.6666667 1.3333333 FALSE
```

We can see in the table that no points have changed groups, and the plot looks the same as the plot in the last iteration. So we stop here, because the algorithm has converged.

Question 2

Suppose DNA bases in a protein-coding region follow the distribution:

DNA base	Probability
A	θ
C	$\frac{1}{4}$
G	$\frac{1}{2}$
T	$\frac{1}{4} - \theta$

In an experiment, the number of observed “A” or “C” bases at the position is x , and the number of observed “G” or “T” bases at the position is y . The EM algorithm can be used to find parameter θ . Describe the Expectation step and Maximization step in the EM algorithm (25 points)

Answer 2

If the number of A, C, G, and T bases are a , c , g , and t respectively, then the likelihood is

$$P(a, c, g, t | \theta) = C(\theta)^a \left(\frac{1}{4}\right)^c \left(\frac{1}{2}\right)^g \left(\frac{1}{4} - \theta\right)^t$$

We take the log, take the derivative with respect to θ , and set it equal to 0 to solve for θ .

$$\ln P(a, c, g, t | \theta) = \ln C + a \ln \theta + c \ln \frac{1}{4} + g \ln \frac{1}{2} + t \ln \left(\frac{1}{4} - \theta\right)$$

$$\frac{\partial}{\partial \theta} \ln P(a, c, g, t | \theta) = \frac{a}{\theta} - \frac{t}{\frac{1}{4} - \theta} = 0$$

$$\theta = \frac{a}{4(t + a)}$$

We know $\frac{a}{c} = \frac{\theta}{\frac{1}{4}}$ and $a + c = x$, so $a = \frac{c\theta}{\frac{1}{4}} = \frac{(x-a)\theta}{\frac{1}{4}} = \frac{\theta}{\frac{1}{4} + \theta}x$. We also know $\frac{t}{g} = \frac{\frac{1}{4} - \theta}{\frac{1}{2}}$ and $g + t = y$, so $t = \frac{(\frac{1}{4} - \theta)g}{\frac{1}{2}} = \frac{(\frac{1}{4} - \theta)(y - t)}{\frac{1}{2}} = \frac{\frac{1}{4} - \theta}{\frac{3}{4} - \theta}y$. For the algorithm, we give an initial value to θ , and if we know the value of x and y , we solve for a and t . Using these initial values we solve for a new θ ,

$$\theta_{new} = \frac{a}{4(t + a)}$$

We then use θ_{new} to find a_{new} and t_{new} . The steps are repeated until convergence.
 Suppose $x = 12,000$ and $y = 13,000$ and the initial θ is 0.1. Below is the EM algorithm.

```
theta <- 0.1
calcA <- function(theta, x = 12000){
  (x * theta)/((1/4) + theta)
}
calcT <- function(theta, y = 13000){
  (((1/4) - theta) * y)/((3/4)-theta)
}
calcTheta <- function(a, t){
  a/(4 * (t+ a))
}

em.df <- data.frame(n = 0, a = NA, t = NA, theta = theta)
theta.diff <- theta
while(theta.diff > 0.00001){
  n.last = em.df[nrow(em.df), "n"]
  theta.old = em.df[em.df$n == n.last, "theta"]
  a.new = calcA(theta.old)
  t.new = calcT(theta.old)
  theta.new = calcTheta(a.new, t.new)
  n.new = n.last + 1
  em.new.df = data.frame(n = n.new, a = a.new,
                        t = t.new, theta = theta.new)
  em.df <-- rbind(em.df, em.new.df)
  theta.diff <- abs(theta.new - theta.old)
  if(n.new > 1000) break
}
em.df
```

##	n	a	t	theta
## 1	0	NA	NA	0.1000000
## 2	1	3428.571	3000.0000	0.1333333
## 3	2	4173.913	2459.4595	0.1573074
## 4	3	4634.555	2033.1017	0.1737700
## 5	4	4920.688	1719.7820	0.1852538
## 6	5	5107.469	1490.4059	0.1935270
## 7	6	5236.038	1319.2885	0.1996864
## 8	7	5328.684	1188.5530	0.2044073
## 9	8	5397.994	1086.3503	0.2081164
## 10	9	5451.446	1004.8036	0.2110918
## 11	10	5493.704	938.5760	0.2135209
## 12	11	5527.799	883.9650	0.2155335
## 13	12	5555.780	838.3403	0.2172222
## 14	13	5579.073	799.7915	0.2186546
## 15	14	5598.698	766.9020	0.2198810
## 16	15	5615.405	738.6021	0.2209395
## 17	16	5629.755	714.0702	0.2218596
## 18	17	5642.177	692.6657	0.2226644
## 19	18	5653.003	673.8825	0.2233723
## 20	19	5662.493	657.3155	0.2239978
## 21	20	5670.857	642.6370	0.2245530

```

## 22 21 5678.262 629.5793 0.2250478
## 23 22 5684.845 617.9214 0.2254901
## 24 23 5690.720 607.4796 0.2258868
## 25 24 5695.980 598.0996 0.2262436
## 26 25 5700.702 589.6513 0.2265653
## 27 26 5704.954 582.0242 0.2268560
## 28 27 5708.792 575.1235 0.2271192
## 29 28 5712.263 568.8680 0.2273581
## 30 29 5715.409 563.1871 0.2275751
## 31 30 5718.266 558.0198 0.2277727
## 32 31 5720.863 553.3128 0.2279528
## 33 32 5723.229 549.0193 0.2281171
## 34 33 5725.387 545.0981 0.2282673
## 35 34 5727.357 541.5129 0.2284047
## 36 35 5729.158 538.2317 0.2285305
## 37 36 5730.807 535.2257 0.2286458
## 38 37 5732.317 532.4695 0.2287515
## 39 38 5733.701 529.9404 0.2288485
## 40 39 5734.971 527.6179 0.2289377
## 41 40 5736.137 525.4838 0.2290197
## 42 41 5737.209 523.5216 0.2290950
## 43 42 5738.194 521.7163 0.2291644
## 44 43 5739.101 520.0547 0.2292282
## 45 44 5739.935 518.5244 0.2292871
## 46 45 5740.703 517.1145 0.2293413
## 47 46 5741.411 515.8151 0.2293912
## 48 47 5742.063 514.6170 0.2294373
## 49 48 5742.665 513.5118 0.2294798
## 50 49 5743.219 512.4922 0.2295190
## 51 50 5743.731 511.5511 0.2295552
## 52 51 5744.203 510.6823 0.2295887
## 53 52 5744.639 509.8801 0.2296195
## 54 53 5745.042 509.1392 0.2296480
## 55 54 5745.414 508.4546 0.2296744
## 56 55 5745.757 507.8221 0.2296987
## 57 56 5746.075 507.2376 0.2297212
## 58 57 5746.368 506.6973 0.2297420
## 59 58 5746.639 506.1978 0.2297613
## 60 59 5746.890 505.7359 0.2297791
## 61 60 5747.122 505.3088 0.2297955
## 62 61 5747.336 504.9138 0.2298107
## 63 62 5747.534 504.5484 0.2298248
## 64 63 5747.718 504.2105 0.2298378
## 65 64 5747.887 503.8978 0.2298498
## 66 65 5748.044 503.6085 0.2298610
## 67 66 5748.189 503.3408 0.2298713
## 68 67 5748.324 503.0931 0.2298808

```

Question 3

Use the ID3 method to construct a decision tree using the following data set for credit card application. (25 points)

Age	Income	Gender	Risk
<25	>50K	M	High
<25	>50K	F	High
≥ 25	<50K	F	High
≥ 25	>50K	F	Low
≥ 25	>50K	M	Low
<25	<50K	M	High

Answer 3

Question 4

A linear SVM is used to analyze a 2-dimensional data set, and the solution α to the quadratic programming program contains only three non-zero elements $\alpha_1 = 0.05$, $\alpha_2 = 0.05$, and $\alpha_3 = 0.1$. The corresponding input data points are $\mathbf{x}_1 = [0, 3]^T$, $y_1 = 1$, $\mathbf{x}_2 = [2, 4]^T$, $y_2 = 1$, $\mathbf{x}_3 = [3, -0.5]^T$, $y_3 = -1$. Compute the weight vector \mathbf{w} and parameter b for the separating line with the maximum margin (15 points). Compute the margin of the separating line (10 points).